

Peterson's Algorithm

- Peterson's algorithm, also called the tie-breaker algorithm
 - It keeps track of which processes are in their critical sections
 - If two processes try to enter then the a deadlock results
 - This is resolved by letting the last one to start the entry protocol proceed

More Mutual Exclusion

G52CON - Concepts of Concurrency

Peterson's Algorithm

```
// Shared data
boolean in1 = false;
boolean in2 = false;
int last = 1;

Process1 {
    initialise_1();
    while (true) {
        in1 = true;
        last = 1;
        while ( in2 &&
                (last==1)) {};
        critical_1();
        in1=false;
        non_critical_1();
    }
}

Process2 {
    initialise_2();
    while (true) {
        in2 = true;
        last = 2;
        while ( in1 &&
                (last==2)) {};
        critical_2();
        in2=false;
        non_critical_2();
    }
}
```

More Mutual Exclusion

G52CON - Concepts of Concurrency

Peterson's Algorithm

- To see how this works we consider the possible situations when Process1 enters its critical section
 - The cases are symmetric for Process2
 - We consider possible stages that Process2 could be at in its protocol as Process1 exists the while loop and enters its critical section

More Mutual Exclusion

G52CON - Concepts of Concurrency

Peterson's Algorithm

- If Process2 is outside the critical section and exclusion protocol
 - in2 is false, so it is safe for Process1 to proceed
 - If Process2 reaches its critical section before Process1 finishes, then in1 will be true, and last will be 2, so it will wait

```
Process2 {
    initialise_2();
    while (true) {
        in2 = true;
        last = 2;
        while ( in1 &&
                (last==2)) {};
        critical_2();
        in2=false;
        non_critical_2();
    }
}
```

More Mutual Exclusion

G52CON - Concepts of Concurrency

Peterson's Algorithm

- If Process2 has just set in2 to be true
 - in2 is true
 - last will be 1, since Process1 has just set it, and Process2 hasn't yet
 - As a result, Process1 will wait in a spin-lock until...

```
Process2 {
    initialise_2();
    while (true) {
        in2 = true;
        last = 2;
        while ( in1 &&
                (last==2)) {};
        critical_2();
        in2=false;
        non_critical_2();
    }
}
```

More Mutual Exclusion

G52CON - Concepts of Concurrency

Peterson's Algorithm

- If Process2 has just set last to be 2
 - last is 2, so it is safe for Process1 to proceed
 - If Process2 reaches its critical section before Process1 finishes, then in1 will be true, and last will still be 2, so it will wait

```
Process2 {
    initialise_2();
    while (true) {
        in2 = true;
        last = 2;
        while ( in1 &&
                (last==2)) {};
        critical_2();
        in2=false;
        non_critical_2();
    }
}
```

More Mutual Exclusion

G52CON - Concepts of Concurrency

Peterson's Algorithm

- If Process2 evaluating its while loop

- in1 and in2 are both true
- last is either 1 or 2, depending on how we got here
- If last is 1, Process2 will proceed, and Process1 will have to wait
- If last is 2, Process1 will proceed, and Process2 will have to wait

```
Process2 {
    initialise_2();
    while (true) {
        in2 = true;
        last = 2;
        while ( in1 &&
                (last==2)) {};
        critical_2();
        in2=false;
        non_critical_2();
    }
}
```

More Mutual Exclusion

G52CON – Concepts of Concurrency

Peterson's Algorithm

- If Process2 is in its critical section

- in1 and in2 are both true
- Process2 must have seen last == 1 to enter the loop, and hasn't changed it
- Process1 only ever sets last to 1
- So last is still 1, and so Process1 waits

```
Process2 {
    initialise_2();
    while (true) {
        in2 = true;
        last = 2;
        while ( in1 &&
                (last==2)) {};
        critical_2();
        in2=false;
        non_critical_2();
    }
}
```

More Mutual Exclusion

G52CON – Concepts of Concurrency

Peterson's Algorithm

- Does this satisfy
 - Mutual exclusion?
 - Yes
 - Absence of livelock?
 - No
 - Absence of unnecessary delay?
 - No
 - Eventual entry?
 - Yes, if scheduler is *weakly fair*

More Mutual Exclusion

G52CON – Concepts of Concurrency

Peterson's Algorithm

- Note that fine detail is important

- If we swapped the two statements shown, then it doesn't work

```
Process1 {
    initialise_1();
    while (true) {
        in1 = true;
        last = 1;
        while ( in2 &&
                (last==1)) {};
        critical_1();
        in1=false;
        non_critical_1();
    }
}
```

More Mutual Exclusion

G52CON – Concepts of Concurrency

Multiple Processes

- Peterson's Algorithm can be extended to more than 2 processes
 - For n processes, each protocol loops through $n-1$ instances of the 2 process protocol
 - If a process makes it through all $n-1$ checks, then it is safe to proceed
 - This is complicated, but can be made to work

More Mutual Exclusion

G52CON – Concepts of Concurrency

Peterson's vs Test-and-Set

- | | |
|--|--|
| <ul style="list-style-type: none"> • Test-and-Set <ul style="list-style-type: none"> • Provides mutual exclusion, absence of deadlock, and of unnecessary delay • Provides eventual entry if scheduler is <i>strongly fair</i> • Requires special operations • Easy to do for >2 processes | <ul style="list-style-type: none"> • Peterson's <ul style="list-style-type: none"> • Provides mutual exclusion, absence of deadlock, and of unnecessary delay • Provides eventual entry if scheduler is <i>weakly fair</i> • Requires no special operations • Difficult to do for >2 processes |
|--|--|

More Mutual Exclusion

G52CON – Concepts of Concurrency

The Ticket Algorithm

- Based on the idea of taking a ticket number and waiting at bakeries etc.
- Each process on entry gets a number from a counter
- They then wait for their number to come up before entering their critical sections

More Mutual Exclusion

G52CON – Concepts of Concurrency

The Ticket Algorithm

- In this approach
- There are two shared variables, **counter** and **next**
- A process gets the value of the counter as its **turn**, then the counter is updated
- Once the process's **turn** is equal to **next** it can go

```
// Process x
initialise_x();
while (true) {
    turn = counter;
    counter++;
    while(turn!=next){};
    critical_x();
    next++;
    non_critical_x();
}
```

More Mutual Exclusion

G52CON – Concepts of Concurrency

The Ticket Algorithm

- There is a problem
- Two processes could interfere with each other and get the same turn
- Can be overcome if we have a fetch-and-add instruction
- FA(x) returns the current value of x then does x++ atomically

```
// Process x
initialise_x();
while (true) {
    turn = FA(counter);
    while(turn!=next){};
    critical_x();
    next++;
    non_critical_x();
}
```

More Mutual Exclusion

G52CON – Concepts of Concurrency

The Bakery Algorithm

- The ticket algorithm needs a special instruction to work
- However, a variant of it, the Bakery algorithm does not
- Uses ideas from Peterson's algorithm to deal with two processes having the same number
- We need to define a comparison between ordered pairs, $(a,b) < (c,d)$, as $(a < c) \vee ((a = c) \wedge (b < d))$

More Mutual Exclusion

G52CON – Concepts of Concurrency

The Bakery Algorithm

```
// For process i out of n total processes
while (true) {
    choosing[i] = true;
    number[i] = max(number[1]...number[n]) + 1;
    choosing[i] = false;
    for j = 1 to n except i {
        while(choosing[j]) {}
        while((number[j] != 0) &&
            ((number[j],j) < (number[i],i)) {}
    }
    critical_i();
    number[i] = 0;
    non_critical_i();
}
```

More Mutual Exclusion

G52CON – Concepts of Concurrency

The Bakery Algorithm

- How it works:
 - Each process has an identifier, **i**
 - To enter a critical section, each process takes a number bigger than all others assigned
 - While taking a number, the process is in the state 'choosing'
 - It then looks at all the other Processes, **j**
 - If they are choosing, it waits until they finish
 - Otherwise, it checks if their numbers are lower
 - In a tie, the process with the lower identifier goes first

More Mutual Exclusion

G52CON – Concepts of Concurrency

The Bakery Algorithm

- This algorithm
 - Meets the requirements of a mutual exclusion protocol
 - Is safe even on multiple processors
 - Turns out to be safe, even if reads and writes to the shared variables interfere with each other – not even these need be atomic actions!

More Mutual Exclusion

G52CON – Concepts of Concurrency

Java

- When a Java program has multiple threads their statements can be interleaved, but
 - They can also be reordered by various things (next slide)
 - (Copies of) values can be kept in caches and registers while being used
 - These things can be done as long as the 'meaning' is unchanged

More Mutual Exclusion

G52CON – Concepts of Concurrency


Ordering of Instructions

- The order of instructions can change
 - By the Java compiler in order to increase efficiency
 - By the CPU when it gets the machine instructions to execute
 - By the memory system when scheduling reads and writes to memory

More Mutual Exclusion

G52CON – Concepts of Concurrency

Java Memory Model



More Mutual Exclusion

G52CON – Concepts of Concurrency

Java Memory Model

- Each thread has a *working memory*
 - This is an abstraction of the registers and caches on the CPU
 - Values can be stored here while being processed, but they are only copies of the value from main memory until any updates are written
 - This improves efficiency

More Mutual Exclusion

G52CON – Concepts of Concurrency

Java Memory Model

- The Java Memory Model specifies when values must be transferred between working and main memory
 - *Atomicity*, when instructions have indivisible effects
 - *Visibility*, when actions of one thread need to be seen by another
 - *Ordering*, when the actions of a thread can be seen to be out of original order

More Mutual Exclusion

G52CON – Concepts of Concurrency

Atomicity

- Reads or writes to memory locations other than longs and doubles
 - Are guaranteed to return the initial value, or some value written by some thread
 - Not guaranteed to return the last value written by the any thread – that is, you could get an out-of-date value

More Mutual Exclusion

G52CON – Concepts of Concurrency

Visibility

- Without synchronisation, there are no guarantees about when one thread's changes are visible to others
 - When a thread accesses an object it sees a value as defined by atomicity
 - When a thread finishes any written values are flushed to main memory

More Mutual Exclusion

G52CON – Concepts of Concurrency

Ordering

- The order of operations may change
 - From a single thread's point of view they appear to occur in the original order (as-if-serial semantics)
 - From another thread's point of view, almost anything can happen

More Mutual Exclusion

G52CON – Concepts of Concurrency

Synchronisation

- If code is properly **synchronized** then this is simplified
 - The **synchronized** blocks are atomic and visible with respect to any other **synchronized** code with the same lock (usually on the same object)
 - Processing of code in **synchronized** code is in the program specified order

More Mutual Exclusion

G52CON – Concepts of Concurrency

Peterson's Algorithm

- Peterson's algorithm relies on
 - Atomicity – when it reads and writes variables
 - Visibility – so that processes can see what each other are doing when they are in the entry protocol
 - Ordering – if we set **in1** and **last** in the wrong order it doesn't work

More Mutual Exclusion

G52CON – Concepts of Concurrency

Spin Locks in Java

- Without synchronization
 - We might have to wait for an arbitrarily long time for values of **in1**, **in2**, and **last** to become visible
 - The order of setting **in1** and **last** could be changed – the only guarantee is that the thread itself can't tell the difference, and there is no real difference until interference occurs

More Mutual Exclusion

G52CON – Concepts of Concurrency

Java Scheduler Fairness

- Most protocols need the scheduler to be weakly or strongly fair
 - The Java spec makes no guarantees, not even that threads get to make progress
 - Most implementations have some sort of fairness with respect to running threads
 - You can't rely on it though

More Mutual Exclusion

G52CON – Concepts of Concurrency

Spin Locks and Priority

- Java lets you set the priority of threads
 - In general a high-priority thread gets to go before a low-priority thread
 - If a high-priority thread is in a spin lock, waiting for a low-priority thread to exit its critical section then it might spin forever

More Mutual Exclusion

G52CON – Concepts of Concurrency